

Logistic Regression for Classification Problems

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1+e^{-z}} \rightarrow \text{Sigmoid function}$$

$$h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$

$$0 \leq h_{\theta}(x) \leq 1$$

$h_{\theta}(x)$ is going to be conceptually represented as the probability that $y=1$



$$y \in \{0, 1\}$$

Cost Function:

Linear θ

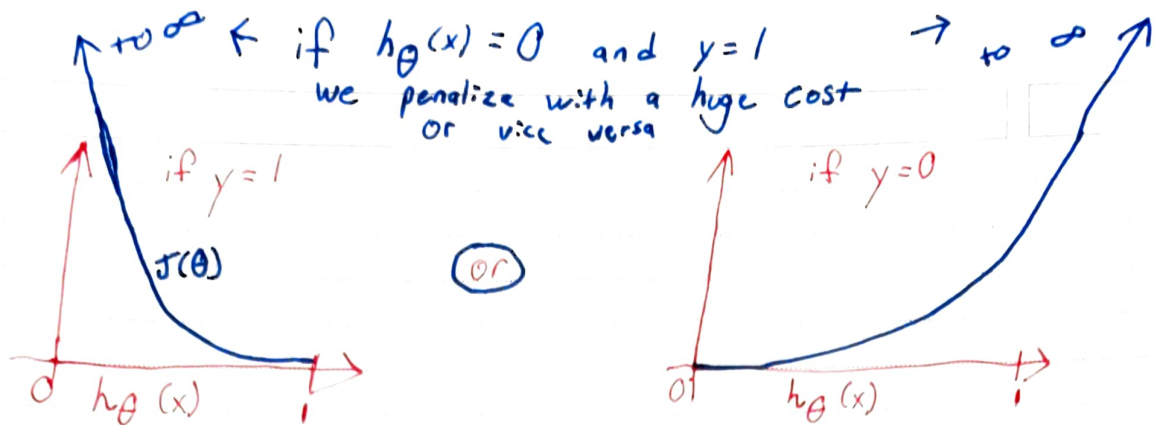
$$\text{cost}(h_{\theta}(x^{(i)}), y) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

if using the linear regression with the sigmoid definition of $h_{\theta}(x)$, would be "non-convex"



Logistic Regression:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$



Logistic Regression Cost Function:

$$J(\theta) = \frac{1}{n} \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

Gradient Descent: Univariate

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$\theta := \theta - \alpha \frac{1}{n} \sum_{i=1}^m [(h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}]$$

$$\theta := \theta - \frac{\alpha}{n} X^T (y(X\theta) - \vec{y})$$

Week 3 - Overfitting

Regularization helps with overfitting by making theta smaller.

Gradient Descent:

$$\theta_j := \theta_j - \alpha \left[\frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{n} \theta_j \right]$$

Cost Function:


$$J(\theta) = \left[-\frac{1}{n} \sum_{i=1}^n y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] + \frac{\lambda}{2n} \sum_{j=1}^n \theta_j^2$$

We can find performance issues ... ie. low cost but not the lowest cost through gradient checking

Gradient Checking:

Week 4 - Neural Networks

$\alpha_0^{(1)} = x$ * can be thought of as the 1st activation layer



Layer (1)
(input)

Layer (2)
(hidden)

Layer (3)
(output)

$a_i^{(j)}$ = activation of unit i in layer j

θ_j = matrix of weights controlling function mapping from Layer j to $j+1$

$$a_1^{(2)} = g(\theta_{10}^{(1)} x_0 + \theta_{11}^{(1)} x_1 + \theta_{12}^{(1)} x_2 + \theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\theta_{20}^{(1)} x_0 + \theta_{21}^{(1)} x_1 + \theta_{22}^{(1)} x_2 + \theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\theta_{30}^{(1)} x_0 + \theta_{31}^{(1)} x_1 + \theta_{32}^{(1)} x_2 + \theta_{33}^{(1)} x_3)$$

$$h_\theta(x) = a_1^{(3)} = g(\theta_{10}^{(2)} a_0^{(2)} + \theta_{11}^{(2)} a_1^{(2)} + \theta_{12}^{(2)} a_2^{(2)} + \theta_{13}^{(2)} a_3^{(2)})$$

Vectorized Approach

Previously $g(z) = g(\theta^T x)$ therefore:

now $z_1^{(2)} = \theta_{10}^{(1)} x_0 + \theta_{11}^{(1)} x_1 + \theta_{12}^{(1)} x_2 + \theta_{13}^{(1)} x_3$

$$z_2^{(2)} = \theta_{20}^{(1)} x_0 + \theta_{21}^{(1)} x_1 + \theta_{22}^{(1)} x_2 + \theta_{23}^{(1)} x_3$$

$$z_3^{(2)} = \theta_{30}^{(1)} x_0 + \theta_{31}^{(1)} x_1 + \theta_{32}^{(1)} x_2 + \theta_{33}^{(1)} x_3$$

$$z^{(2)} = \theta^{(1)} x \text{ or } z^{(2)} = \theta^{(1)} \alpha^{(1)}$$

$$a^{(2)} = g(z^{(2)}) \leftarrow \text{element-wise sigmoid of } z$$

* We also need to add a bias - $\alpha_0^{(2)}$ *

$$z^{(3)} = \theta^{(2)} \alpha^{(2)}$$

$$h_\theta(x) = a^{(3)} = g(z^{(3)})$$

Week 5 - Neural Network Backpropagation

Cost Function

$$J(\theta) = -\frac{1}{n} \sum_{i=1}^m \sum_{k=1}^K \left[y_k^{(i)} \log((h_{\theta}(x^{(i)}))_k) + (1 - y_k^{(i)}) \log(1 - (h_{\theta}(x^{(i)}))_k) \right] + \frac{\lambda}{2m} \sum_{l=1}^{L-1} \sum_{i=1}^{S_l} \sum_{j=1}^{S_{l+1}} (\theta_{ij}^{(l)})^2$$

* not ith training example

L = total # of layers

S_l = # of units (not counting bias) in layer l

K = # of output units / classes

j = node in layer l

$$g'(z^{(3)}) = a^{(3)} * (1 - a^{(3)})$$


To calculate error for any given node:

$$\delta_j^{(4)} = a_j^{(4)} - y_j \quad \text{or} \quad \delta^{(4)} = a^{(4)} - y \rightarrow \delta^{(3)} = (\theta^{(3)})^T \delta^{(4)} * g'(z^{(3)})$$

of output units

NO $\delta^{(1)}$

Gradient Checking: finds "buggy" or inefficient implementation



$$\text{QER: } \frac{d}{d\theta} J(\theta) \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon} \quad \epsilon \approx 10^{-4}$$

Vectorized:

$$\frac{d}{d\theta_1} J(\theta) \approx \frac{J(\theta_1 + \epsilon, \theta_2, \theta_3, \dots, \theta_n) - J(\theta_1 - \epsilon, \theta_2, \theta_3, \dots, \theta_n)}{2\epsilon}$$

$$\frac{d}{d\theta_2} J(\theta) \approx \frac{J(\theta_1, \theta_2 + \epsilon, \theta_3, \dots, \theta_n) - J(\theta_1, \theta_2 - \epsilon, \theta_3, \dots, \theta_n)}{2\epsilon}$$

$$\vdots$$

$$\frac{d}{d\theta_n} J(\theta)$$

continued

Week 5 - Neural Network Backpropagation

Gradient Checking Implementation:

theta is a vector (unrolled) ..

```
for i = (1:n);  
    thetaPlus = theta  
    thetaPlus(i) += Epsilon  
    thetaMinus = theta  
    thetaMinus(i) -= Epsilon  
    gradApprox(i) = (J(thetaPlus) - J(thetaMinus)) / (2 * Epsilon)  
end;
```

1. Implement backprop to compute DVec (unrolled deltas)
2. Implement Gradient Checking to compute gradApprox
3. Turn off gradient checking before training

Random Initialization: How can we initialize theta?

* We cannot initialize all thetas to zero or all neurons will be the same.

aka Symmetry Breaking

Set each θ to a random value between $-\epsilon$ and ϵ

✓ 10×11 matrix between 0 and 1

theta1 = rand(10, 11) * (2 * epsilon) - epsilon*

* not the same as epsilon in gradient checking

Week 5 - Overview for Neural Networks

1. Network Architecture

- input
- output
- hidden (usually 1 layer or same # of units per layer)
 - more units the better (usually more than # of inputs)

2. Training

1. Randomly initialize weights for θ
2. Implement forward propagation to get $h_{\theta}(x^{(i)})$ or a
3. Implement cost function to get $J(\theta)$
4. Implement backward propagation to get $\frac{\partial}{\partial \theta_{j,k}} J(\theta)$
for $i = 1:m$
 - forward prop (); returns a (activations)
 - backward prop (); returns $\frac{\partial}{\partial \theta_{j,k}} J(\theta)$ (derivative)
5. Use gradient checking to verify backprop, then disable it.
6. Use a optimization function like gradient descent to minimize $J(\theta)$

English Translations

1. Where do we start on the hill?
2. What is our guess? hypothesis?
3. How wrong was our guess?
4. Which direction to the bottom of the hill? (compass)
5. Is our compass accurate?
6. Use all the info above to make better guesses...each iteration.

$$\frac{\partial}{\partial \theta_{j,k}} J(\theta) = a_j^{(k)} \delta_i^{(k+1)*} \quad * \text{ignoring } \lambda \text{ if } \lambda = 0 *$$


Week 5 - Mechanical Implementation of Backprop

Forward Prop:

$y_bin = eye(num_labels)[y.T]$ # create an identity matrix for y

$bias = np.ones((m, 1))$ # create a bias vector
 $noBias1 = np.delete(Theta1, 0, 1)$
 $noBias2 = np.delete(Theta2, 0, 2)$

$a1 = concat([bias, X], axis=1)$
 $z2 = a1 @ Theta1.T$
: # work forward



Cost Function:

$$J = (-1/m) * np.sum(y_bin * \log(h) + (1 - y_bin) * \log(1 - h))$$

Regularization:

$$reg = (lambda / (2 * m)) * (np.sum(noBias1 ** 2) + np.sum(noBias2 ** 2))$$

$$J += reg$$

Backprop:

$d3 = a3 - y_bin$
 $d2 = d3 @ noBias2 * sigmoid_gradient(z2)$
 $d1 = \text{"no such thing as error for layer 1"}$

Δ
 δ
 $\Delta_1 = d2.T @ a1$ # use Δ_1 as an "accumulator" for our δ
 $\Delta_2 = d3.T @ a2$

$\Theta1_grad = (1/m) * \Delta_1$ # Adjust our respective theta
 $\Theta2_grad = (1/m) * \Delta_2$

Add our regularization term

$\Theta1_grad[:, 1:] = \Theta1_grad[:, 1:] + 1 * lambda / m * \Theta1[:, 1:]$
 $\Theta2_grad[:, 1:] = \Theta2_grad[:, 1:] + 1 * lambda / m * \Theta2[:, 1:]$

$grad = np.concatenate([\Theta1_grad.ravel(), \Theta2_grad.ravel()])$